

# Web-scale architecture

---

Prajeesh Prathap

# Agenda

- ❑ Web-Scale architecture Introduction
- ❑ Eventual Consistency
- ❑ CQRS
- ❑ Event Sourcing
- ❑ Design for failure
- ❑ Domain Driven Design
- ❑ Onion Architecture
- ❑ Microservices
- ❑ Actor Model

I am going to talk about web-scale architecture and how it allows organizations to build systems in an agile manner that offer high levels of availability and flexibility and support Continuous delivery method



**Praieesh Prathap**  
IT Architect



CONTACT












# Traditional Architecture

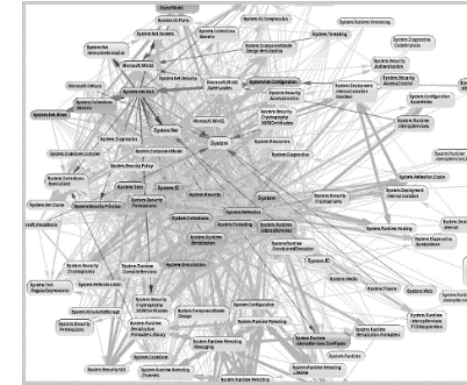
## Monolith



## SOA



## Big Ball of Mud



- **Bad maintainability**

- Lots of tight coupling
- Changes resonate throughout the entire application landscape

- **Bad scalability**

- **Low availability**

- Often offline for upgrades or maintenance
- Services / systems coupled @runtime


- **Long release-cycles**

# Traditional Architecture


How can we change this?



Agile Approach



Continuous Delivery  
Automation  
Dev Ops



Modern Web-scale Architecture



# Introduction of the term web-scale

*In a research note that was published yesterday, Gartner introduced the term "web-scale IT." What is web-scale IT? It's our effort to describe all of the things happening at large cloud services firms such as Google, Amazon, Rackspace, Netflix, Facebook, etc., that enables them to achieve extreme levels of service delivery as compared to many of their enterprise counterparts.*

*In addition, while the term "scale" usually refers to size, we're not suggesting that only large enterprises can benefit. Another scale "attribute" is speed and so we're stating that even smaller firms (or departments within larger IT organizations) can still find benefit to a web-scale IT approach. Agility has no size correlation so even more modestly-sized organizations can achieve some of the capabilities of an Amazon, etc., provided that they are willing to question conventional wisdom where needed.*



A web-scale architecture enables building systems that offer scalability, high performance and high availability. It also encourages loose-coupling which enables teams to employ continuous delivery for the development of the system.

Architecture pattern based on small, specialized and autonomous services that communicate using events. This pattern enables agile teams to develop services autonomously and release frequently.

See also: **The Reactive Manifesto**  
<http://www.reactivemaneifesto.org/>

An approach in which functional domains are split into autonomous areas (bounded contexts) containing one or more domain models (aggregates) which can be accessed through only 1 object (aggregate root). Within a bounded context an ubiquitous language is used to describe the entities and behavior within a domain.

Design pattern that separates the 'write side' (commands) and the 'read side' (queries) of a domain model. For both sides the most appropriate implementation is chosen based on the situation (context).

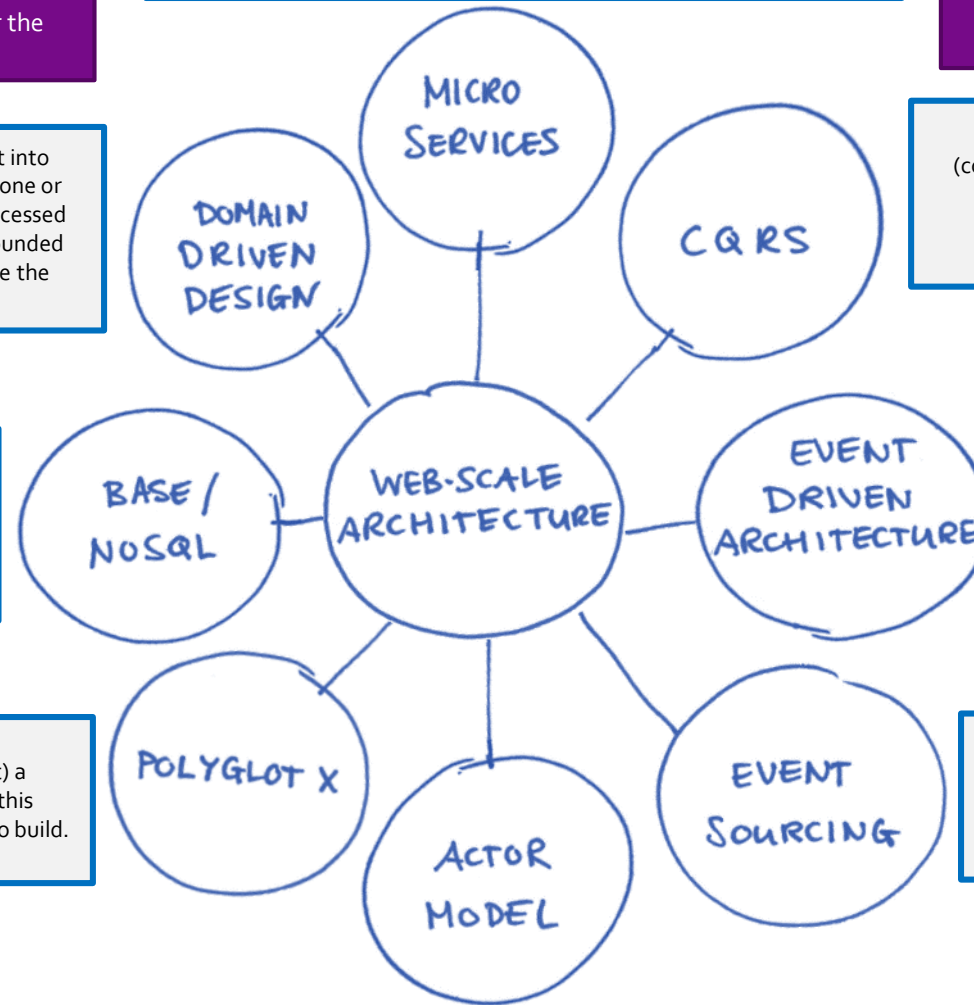
BASE vs. ACID.  
"Not Only SQL". Alternative for data persistence that offers more speed, scalability or lower costs. Different types of NoSQL databases exist, each specialist in a specific type of data: documents, graphs, key-value pairs, wide-columns.

Architecture pattern based on asynchronous communication between components rather than synchronous communication. Duplication of data based on events is common in this kind of architecture. This improves scalability and loose coupling.

An approach in which per situation (context) a technical approach is chosen that best suits this situation and the characteristics of the system to build.

Design pattern that stores the state of a domain object as the series of events that occurred over time as opposed to in a normalized data model. Events are 'replayed' on an object after instantiation to recreate the state of the object.

Design pattern that enables dividing a parallel workload over several small fully autonomous actors. An actor receives messages asynchronously, executes some logic and sends messages. An actor can create new actors to offload work.




# Is that all brand new?

- No!
  - **Actor Model** : 1973 [Carl Hewitt]
  - **CQS** : 1988 [Book van by Bertrand Meyer]]
  - **DDD** : 2003 [Book by Eric Evans]
  - **CQRS** : 2009 [Blog post by Greg Young]
- We see more and more organizations adopting the web-scale architecture patterns

# Disclaimer!

- KISS, common sense and software craftsmanship are still the most important tools of an engineer!
- Choose the best fit-for-purpose solution and architecture style based on complexity and risks!
- Every decision is a trade-off!



Onion architecture Event-based architecture

Domain driven design

Domain modelling DevOps

Micro services

Continuous delivery  
Distributed systems Agile

Polyglot X

CAP

Actor model

CQRS

CQS

Event stroming

Hexagonal architecture Bounded context Circuit breaker

Eventual consistency

BASE Partition tolerance

DDD

Web-scale IT Bulkhead

CAP

Consistency

Serverless computing

Cloud computing distributed computing Aggregate

NOSQL

Event sourcing

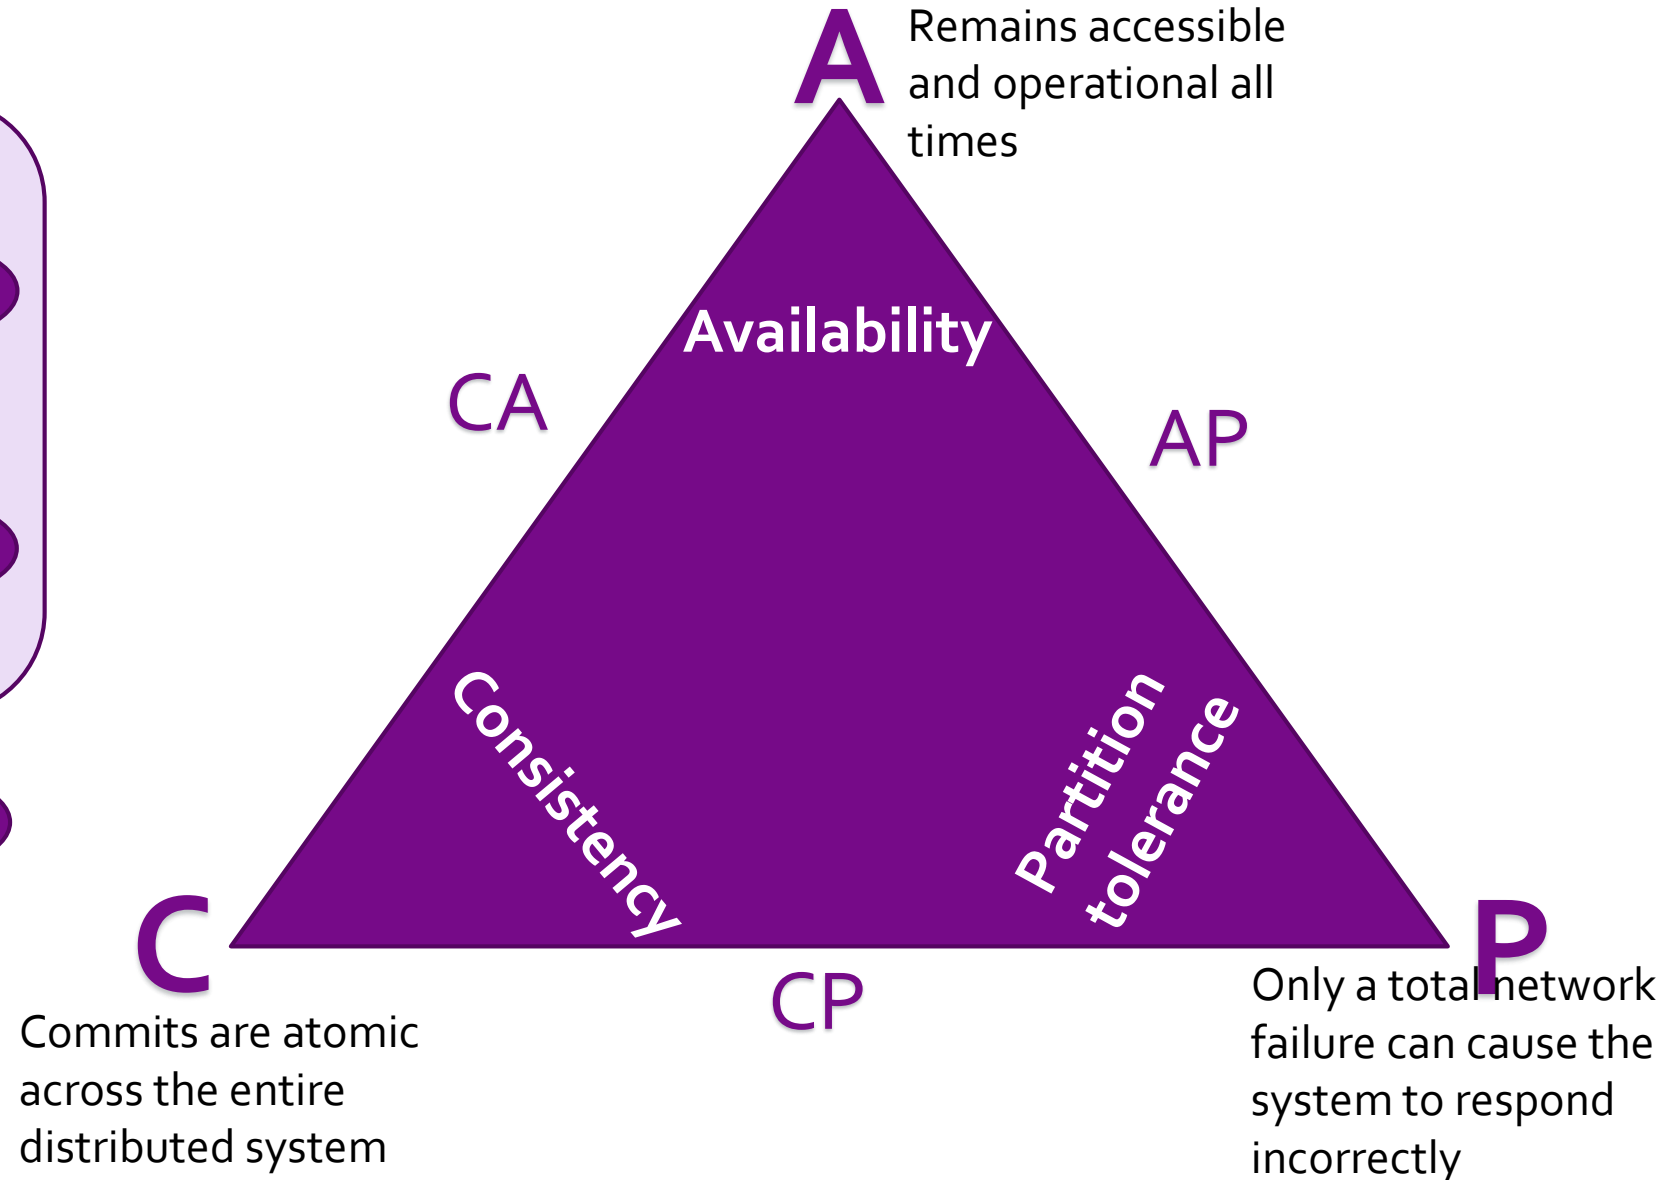

Design for failure

Ubiquitous language

Web-scale architecture




# Eventual consistency – CAP theorem



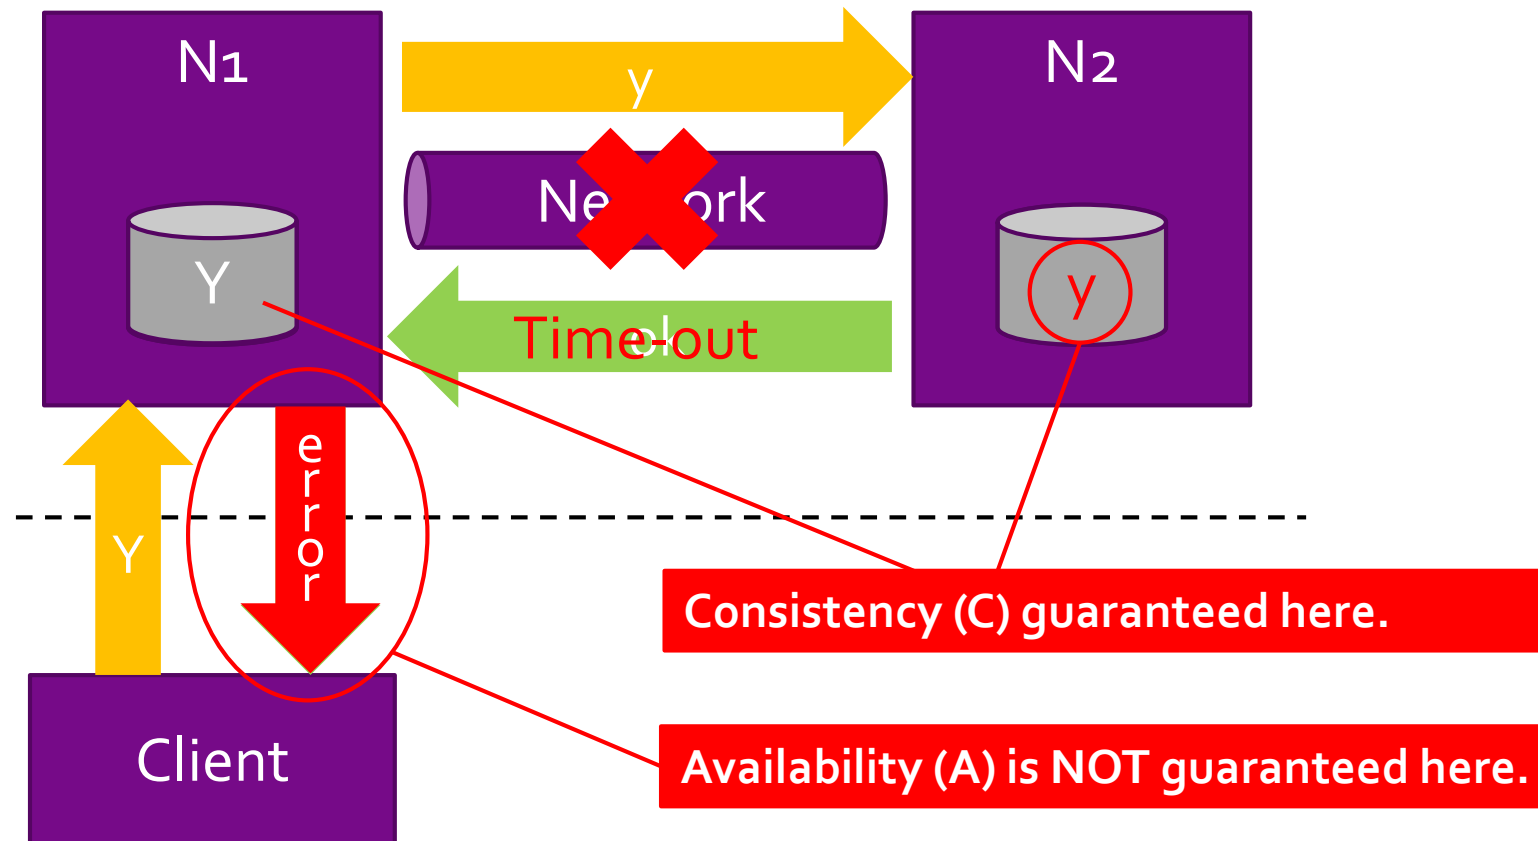
# Eventual consistency - CP

Ensures data is always consistent throughout the system




# Eventual consistency - CP

Ensures data is always consistent throughout the system



# Eventual consistency - AP

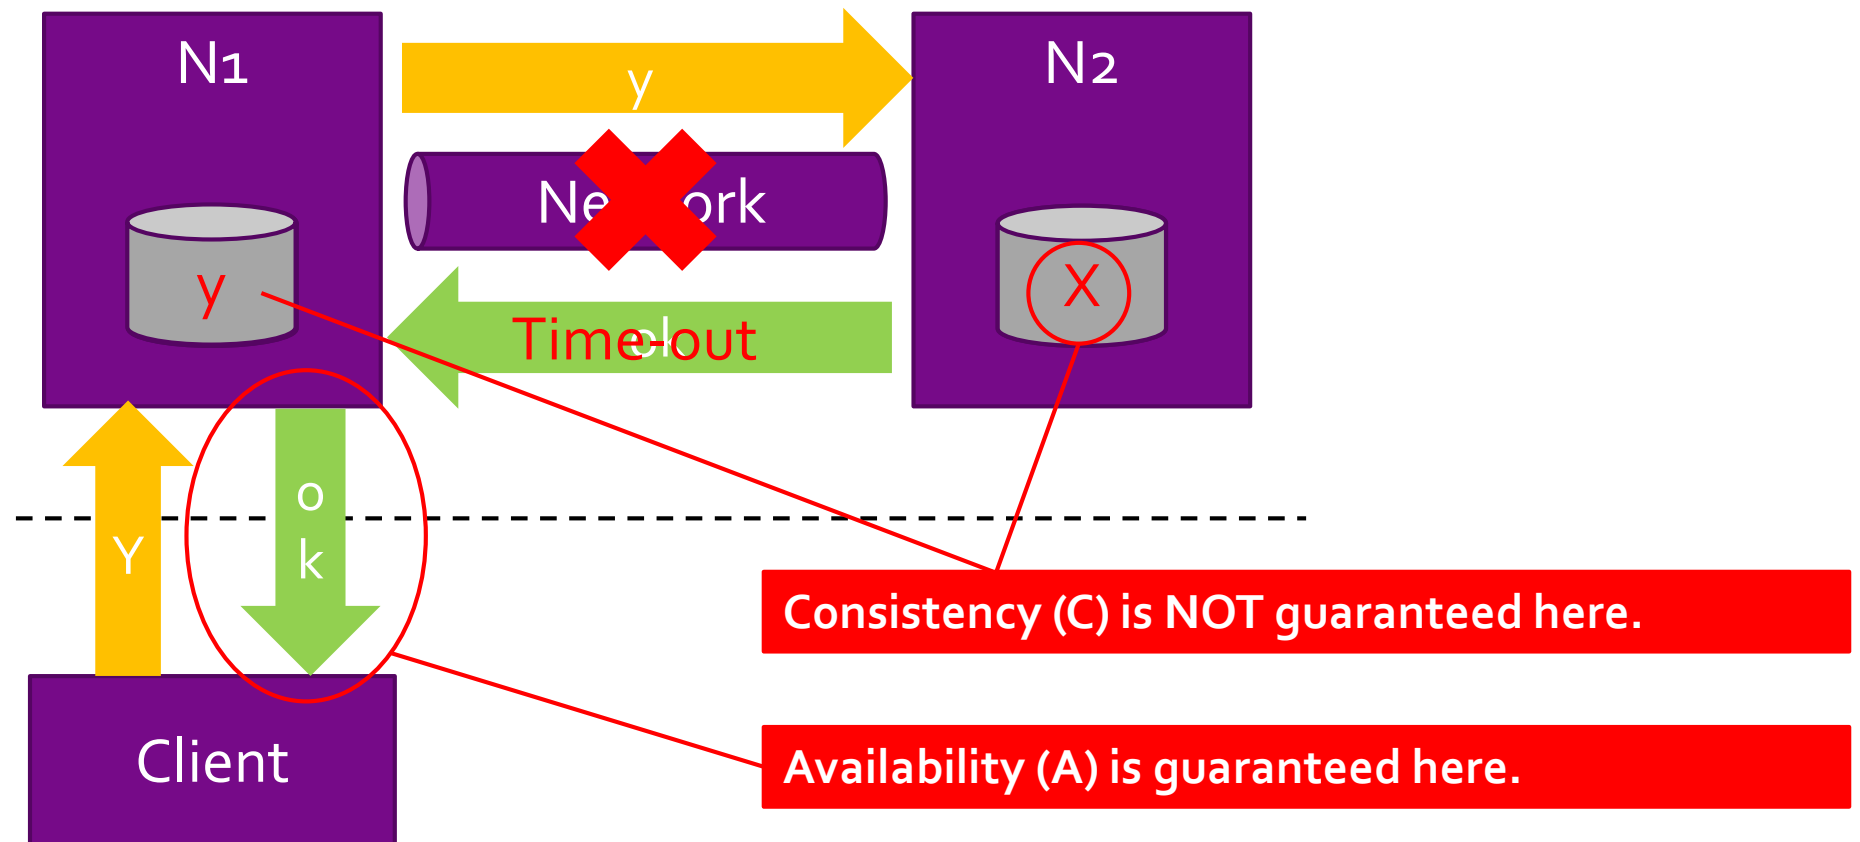
Ensures services are always available





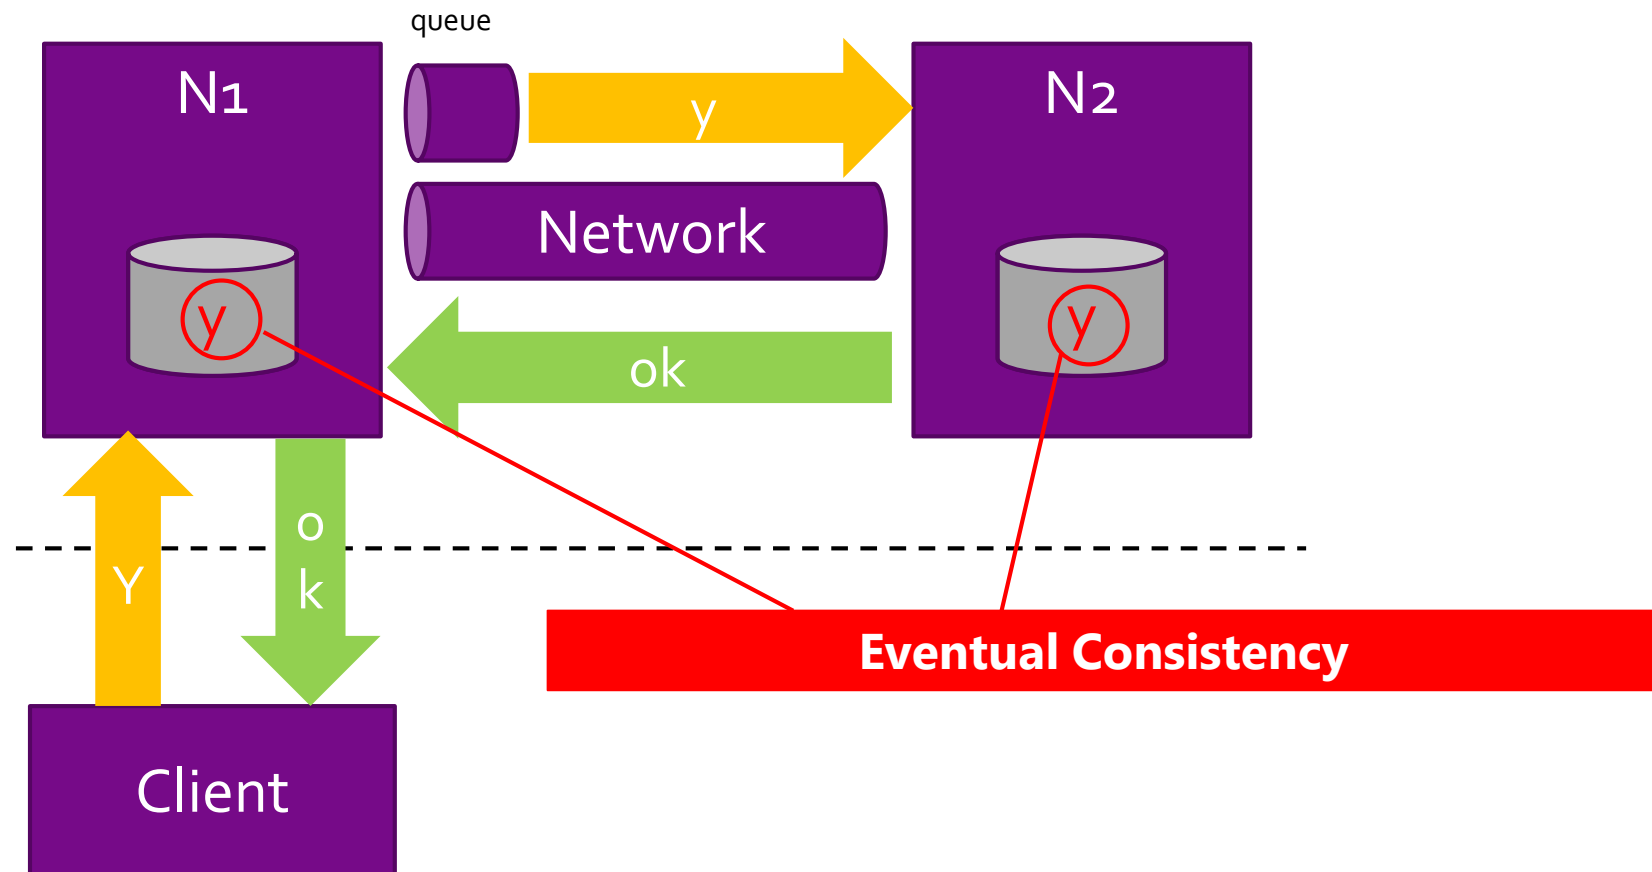
# Eventual consistency - AP

Ensures services are always available



# Eventual consistency - AP

Ensures services are always available



# Eventual consistency

- EC is often not easily accepted
  - “And what about “ACID” and 2PC?”
- Yet, in the “real” world almost every proces is EC
  - Consider whether you really need full consistency when automating business processes
  - Users tend to “get” EC a lot better than we think
  - EC can save you a lot of complexity and trouble (and \$)
  - Compensating actions vs. 2PC

Event-based architecture

Cloud computing

distributed computing Consistency

Eventual consistency

Circuit breaker Hexagonal architecture Agile

Partition tolerance DDD CAP

Domain driven design

Domain modelling DevOps

CQRS

Micro services

Ubiquitous language

Serverless computing

NOSQL Bounded context BASE Actor model

Event stroming

Event sourcing Bulkhead

Aggregate

Distributed systems Web-scale IT Continuous delivery

Design for failure

Polyglot X

Onion architecture




# CQRS

- Command Query Responsibility Segregation
- Pattern that embodies strictly separating updates and queries in a system
  - Scale the update and query parts independently
  - Decreases coupling between systems
  - Enables a task oriented approach for your system (commands)


# Evolution from SOA to CQRS

## Traditional Architecture




# Evolution from SOA to CQRS

CQS




# Evolution from SOA to CQRS

CQRS






# Principle vs Pattern




**Vs.**




# CQRS helps to

*Blocking the user when locking  
the data.*


# CQRS – Airline reservation




# CQRS – Airline reservation




# CQRS – Airline reservation



# CQRS – Airline reservation




# CQRS – Airline reservation





# CQRS – Airline reservation



# Collaborative domain


- Large set of people
- Small set of data
- Locking the data is necessary
- Blocking the user is not

# CQRS – Ticketing agent

Aisle

3 seats  
together

Window

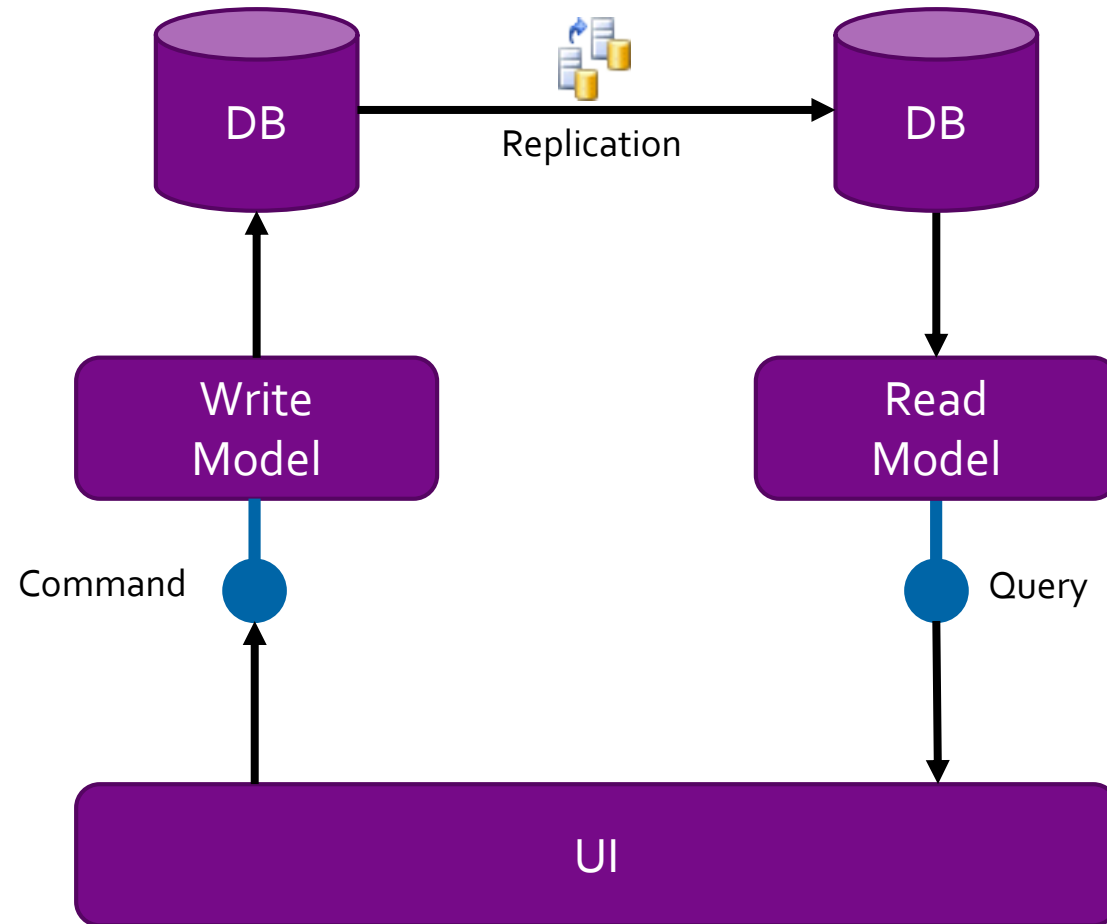


# Tradeoff

- No immediate feedback
- Can't/ don't want to use everywhere.
- Only for collaborative domains.
  - Large number of people
  - Small set of data.

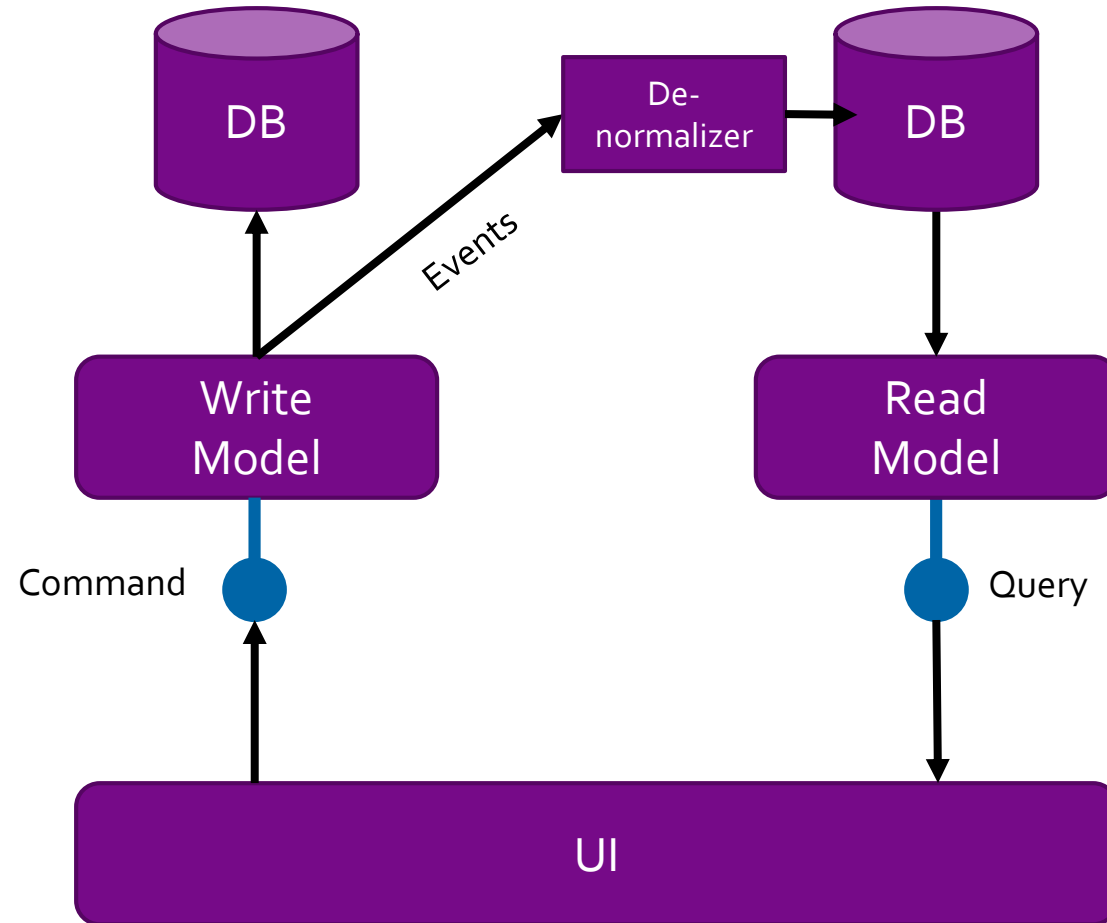
# Evolution from SOA to CQRS

CQRS



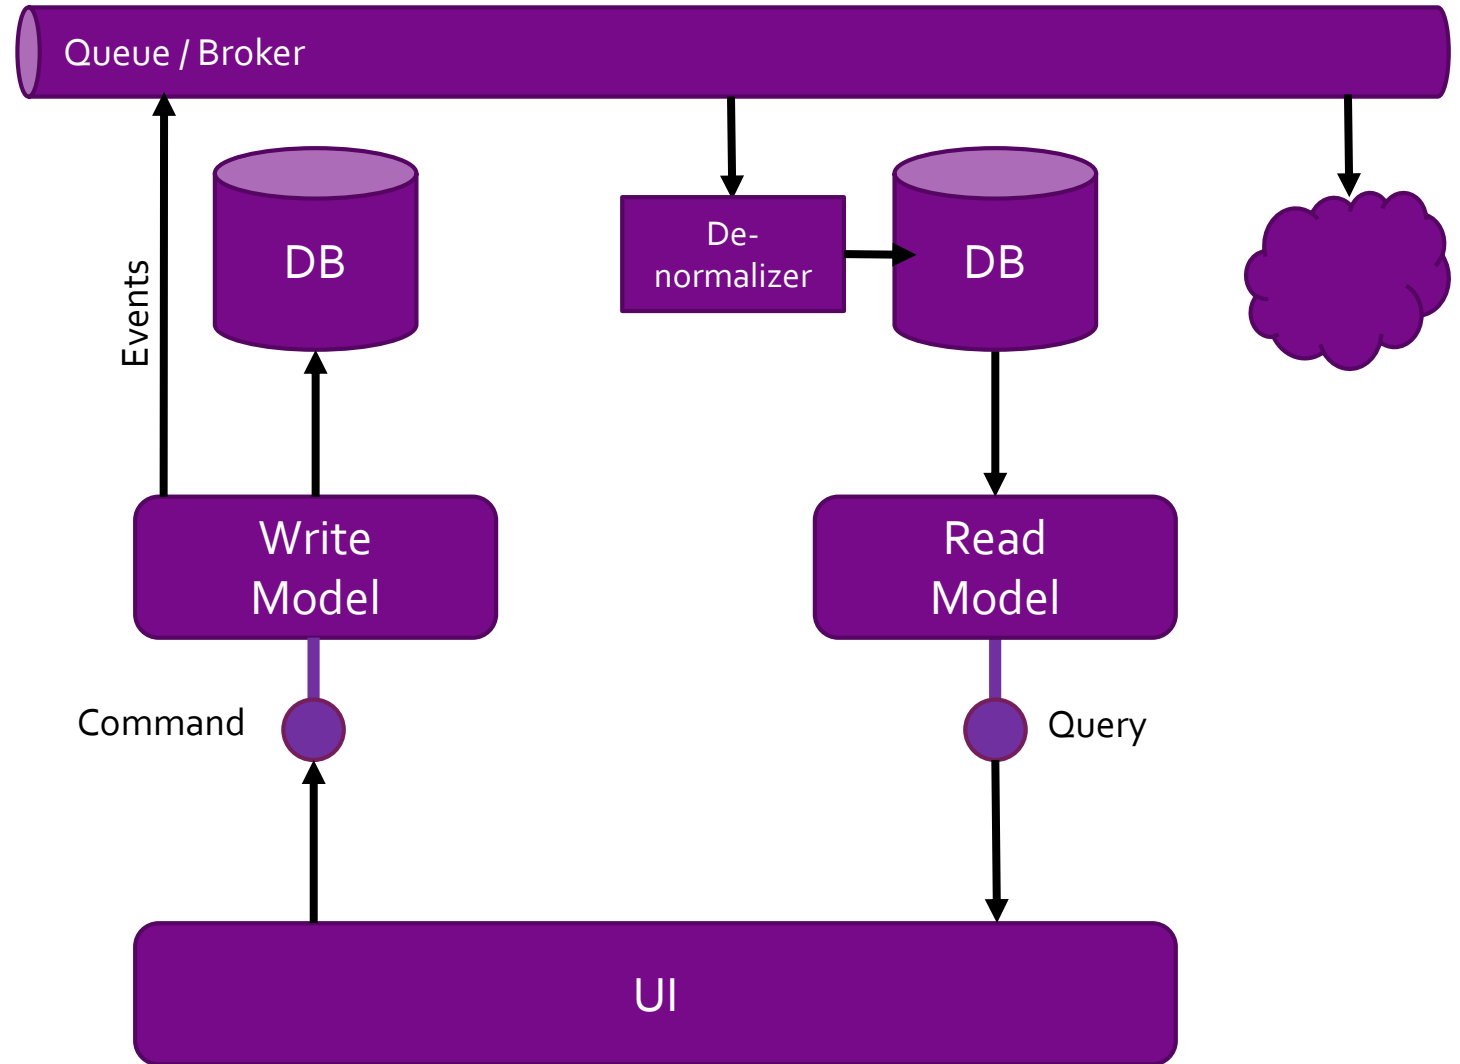
# Evolution from SOA to CQRS

## CQRS




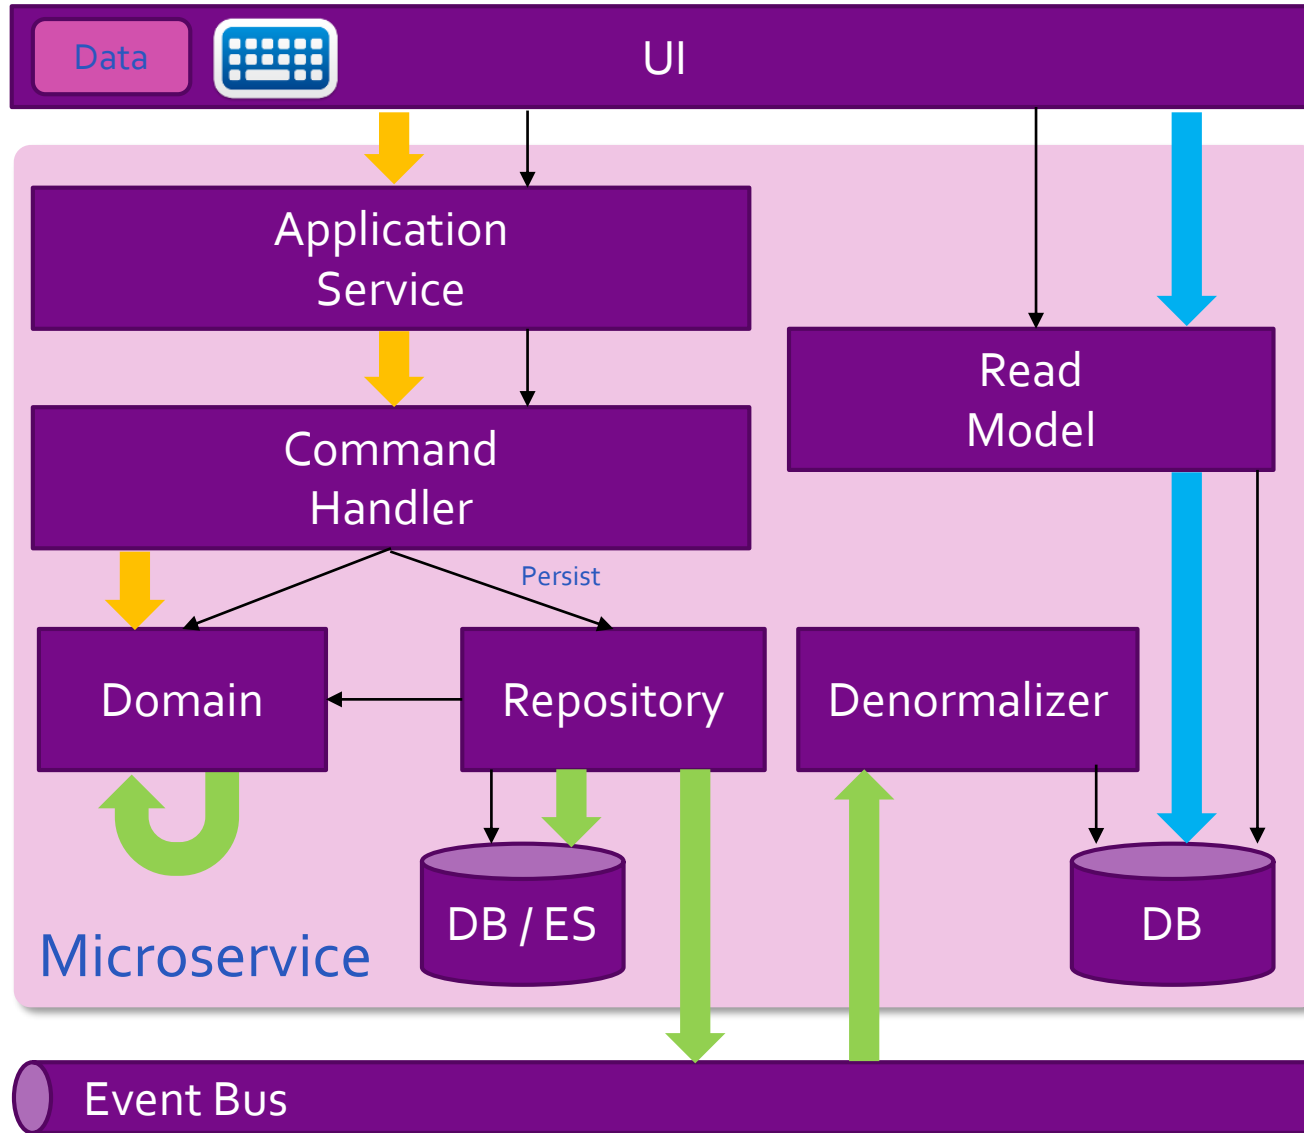
# Evolution from SOA to CQRS

CQRS

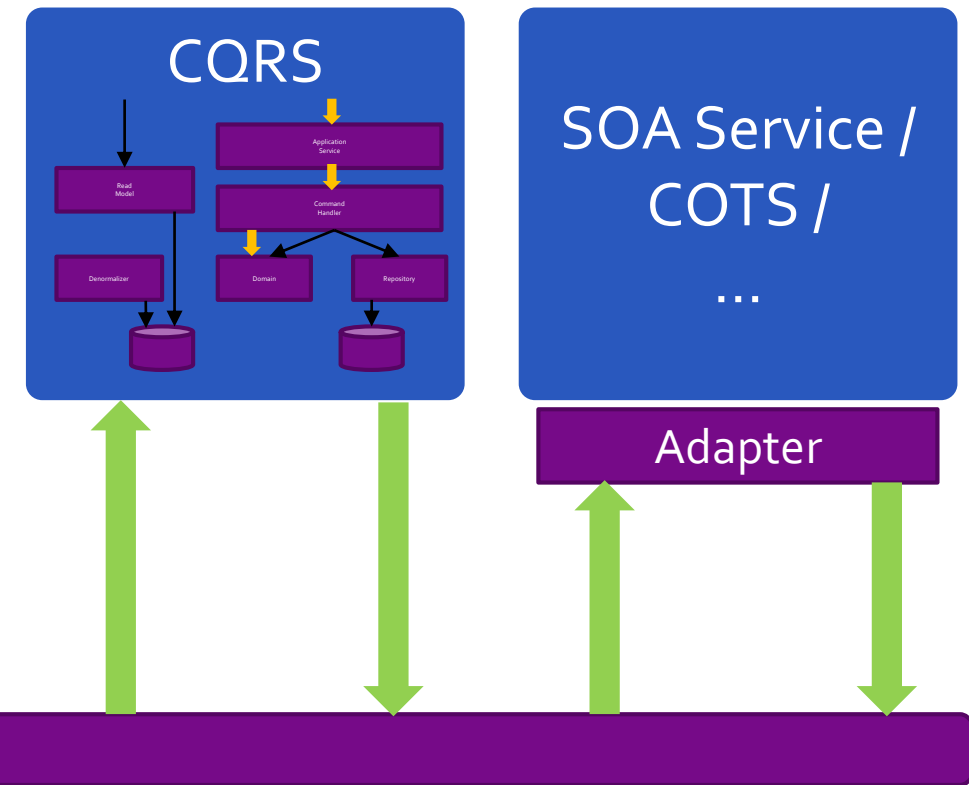




# CQRS - under the covers



## Another domain



# CQRS – CRUD vs. Task based

- Commands are the things that need to be executed
  - Must state business intent
    - So not “UpdateInventory” but “CheckOutItem”
  - Always in the form <Verb><Noun>
  - Can fail (because of business rule / invariants checks)
- Events are things that have happened
  - Always in the form <Noun><Verb (past tense)>
    - CustomerRegistered, ItemCheckedOut, AccountClosed, ...

# CQRS – CRUD vs. Task based

Inventory Item	
Name	<input type="text"/>
Supplier	<input type="text"/>
Count	<input type="text"/>
Status	<div>Status▼</div>
Deactivation comment	<input type="text"/>

# CQRS – CRUD vs. Task based

Name	Supplier	Active
TShirts	GAP	<input checked="" type="checkbox"/>
Keyboards	Logitech	<input checked="" type="checkbox"/>
Mouse	Microsoft	<input checked="" type="checkbox"/>
Laptops	Dell	<input type="checkbox"/>

# CQRS – CRUD vs. Task based

## Deactivate Inventory Item

A comment is required explaining why you are deactivating the inventory item

Cancel

Deactivate

# CQRS

- Handling a command is a 2 phase process:
  - Check phase
    - Check all invariants and business-rules to make sure the command can be executed
    - External resources or services can be called in this phase
  - Execution phase
    - Update the state of the domain
    - No external resources or services can be called in this phase
    - Events are published
- This separation paves the way for *Event Sourcing*

# Web-scale architecture

CQRS

Hexagonal architecture

Aggregate

Serverless computing

Consistency

Micro services

DDD

NOSQL

Onion architecture

Circuit breaker

CAP

Polyglot X

Ubiquitous language

Actor model

CQS

Event sourcing

Bounded context

BASE

CAP

Event stroming

Cloud computing

Bulkhead

Partition tolerance

Distributed systems

DevOps

Domain driven design

Agile

Continuous delivery

Event-based architecture

Web-scale IT

Domain modelling

distributed computing

Design for failure



# Event sourcing



- This is not how the vast majority of applications work today.
- Most applications work by storing the current state and using stored states to process business transactions.

# Event sourcing – structural representation



Ordered list of items



Payment details



Shipping information

# Event sourcing – event representation



Add item 1



Add item 2



Provide payment information



Update item 2



Remove item 1



Provide shipping information

# Event sourcing – business decision

- Is it important to store what was added and removed ✓
- Is it important to store when an item was removed from the cart? ✓

# Event sourcing

- Event-sourcing is an alternative way of persisting the state of your domain-objects
- Not necessarily normalized in an RDBMS, but as a immutable list of events that have occurred over time
- It's about ensuring that all changes made to the application state during the entire lifetime of the application are stored as a sequence of events.

# Event sourcing

- Events are immutable and new events only be appended (not be inserted in between)
  - Think accountant's ledger
  - Appending "Correction" events are allowed
- Snapshots can be used to boost performance
  - Only when absolutely necessary
  - Splitting up the domain can eliminate the need for snapshots

# Event sourcing

- Append only, so super fast (no locking etc.)
- Ability to completely rebuild the state based on event history
- Ability to analyze behavior that occurred in the past
  - Audit log for free
- State can be built-up by issuing events
  - Simplifies automated testing
- Ability to apply changes in retrospect


Design for failure  
Event sourcing Bulkhead  
NOSQL DDD Cloud computing  
Distributed systems Bounded context  
Partition tolerance CAP Web-scale IT  
Actor model CQS Event streaming Circuit breaker  
Polyglot X CQRS BASE CAP Ubiquitous language  
Agile Consistency  
Domain driven design  
Hexagonal architecture Aggregate  
Domain modelling  
Micro services DevOps Serverless computing  
distributed computing Continuous delivery  
Event-based architecture  
Web-scale architecture Onion architecture



# Domain driven design

- Described in Eric Evans' book

*Domain Driven Design – Tackling Complexity in the Heart of Software*



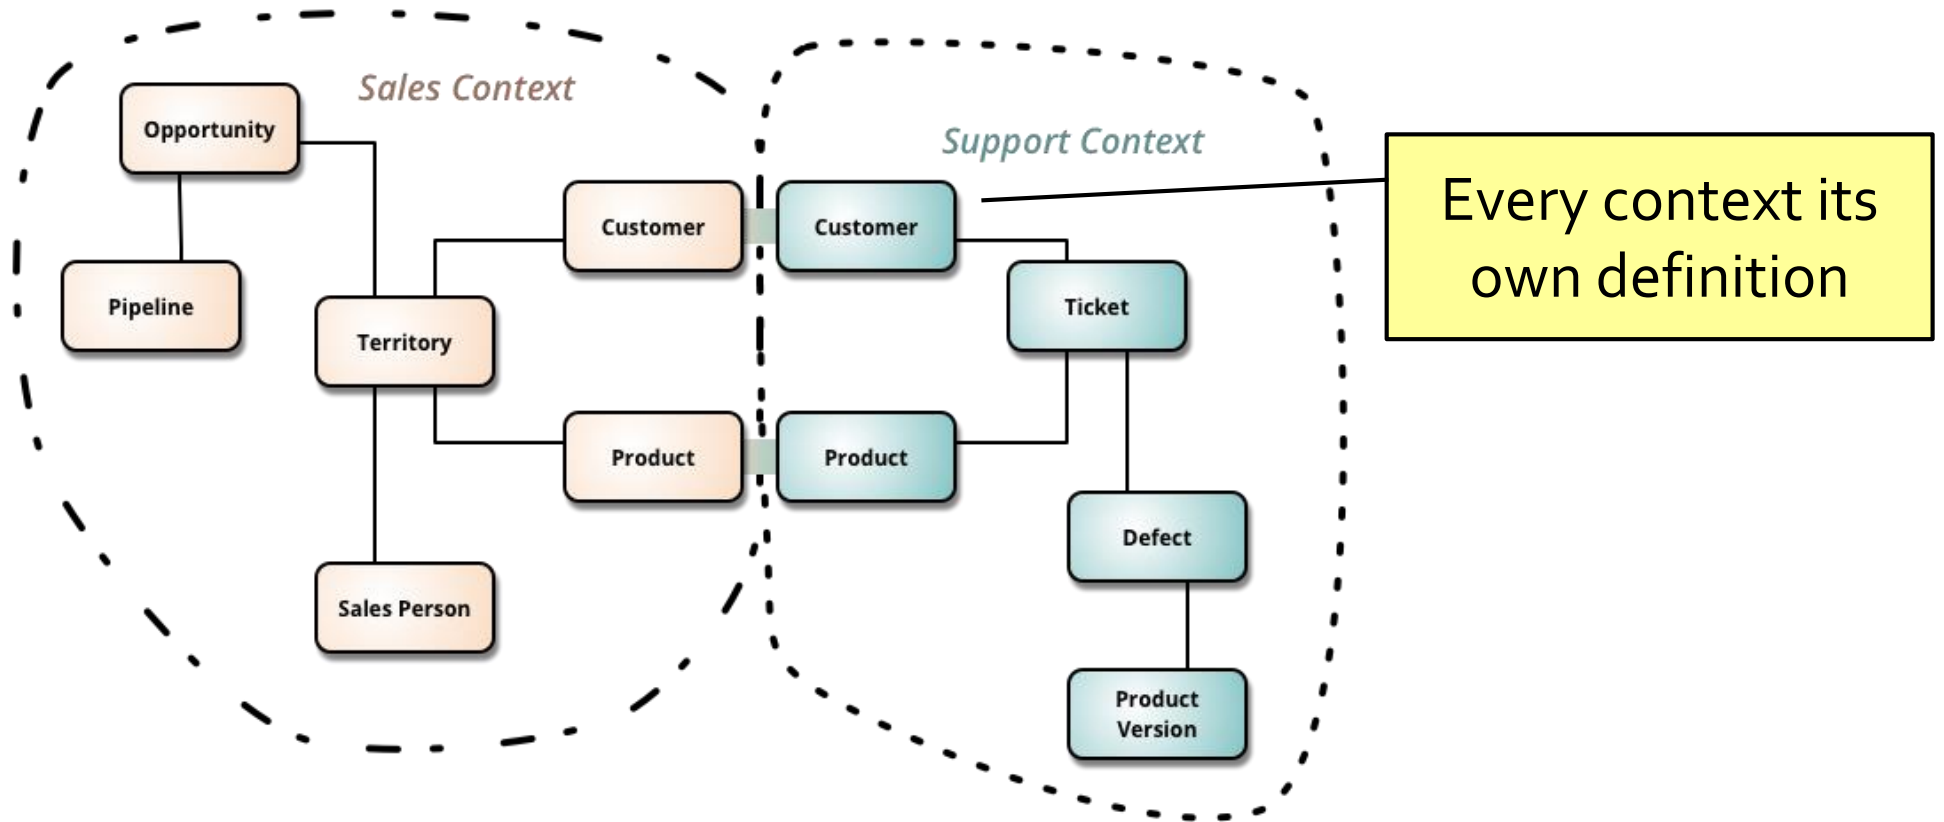
# Domain driven design – Core concepts

- Ubiquitous language
- Bounded contexts
- Aggregate roots

# Domain driven design – Ubiquitous language

- Building up a common, rigorous language between developers and users
- Used by developers and users for common specifications for the domain they are working on.
- Brings all the team members to same page

# Domain driven design – Bounded contexts



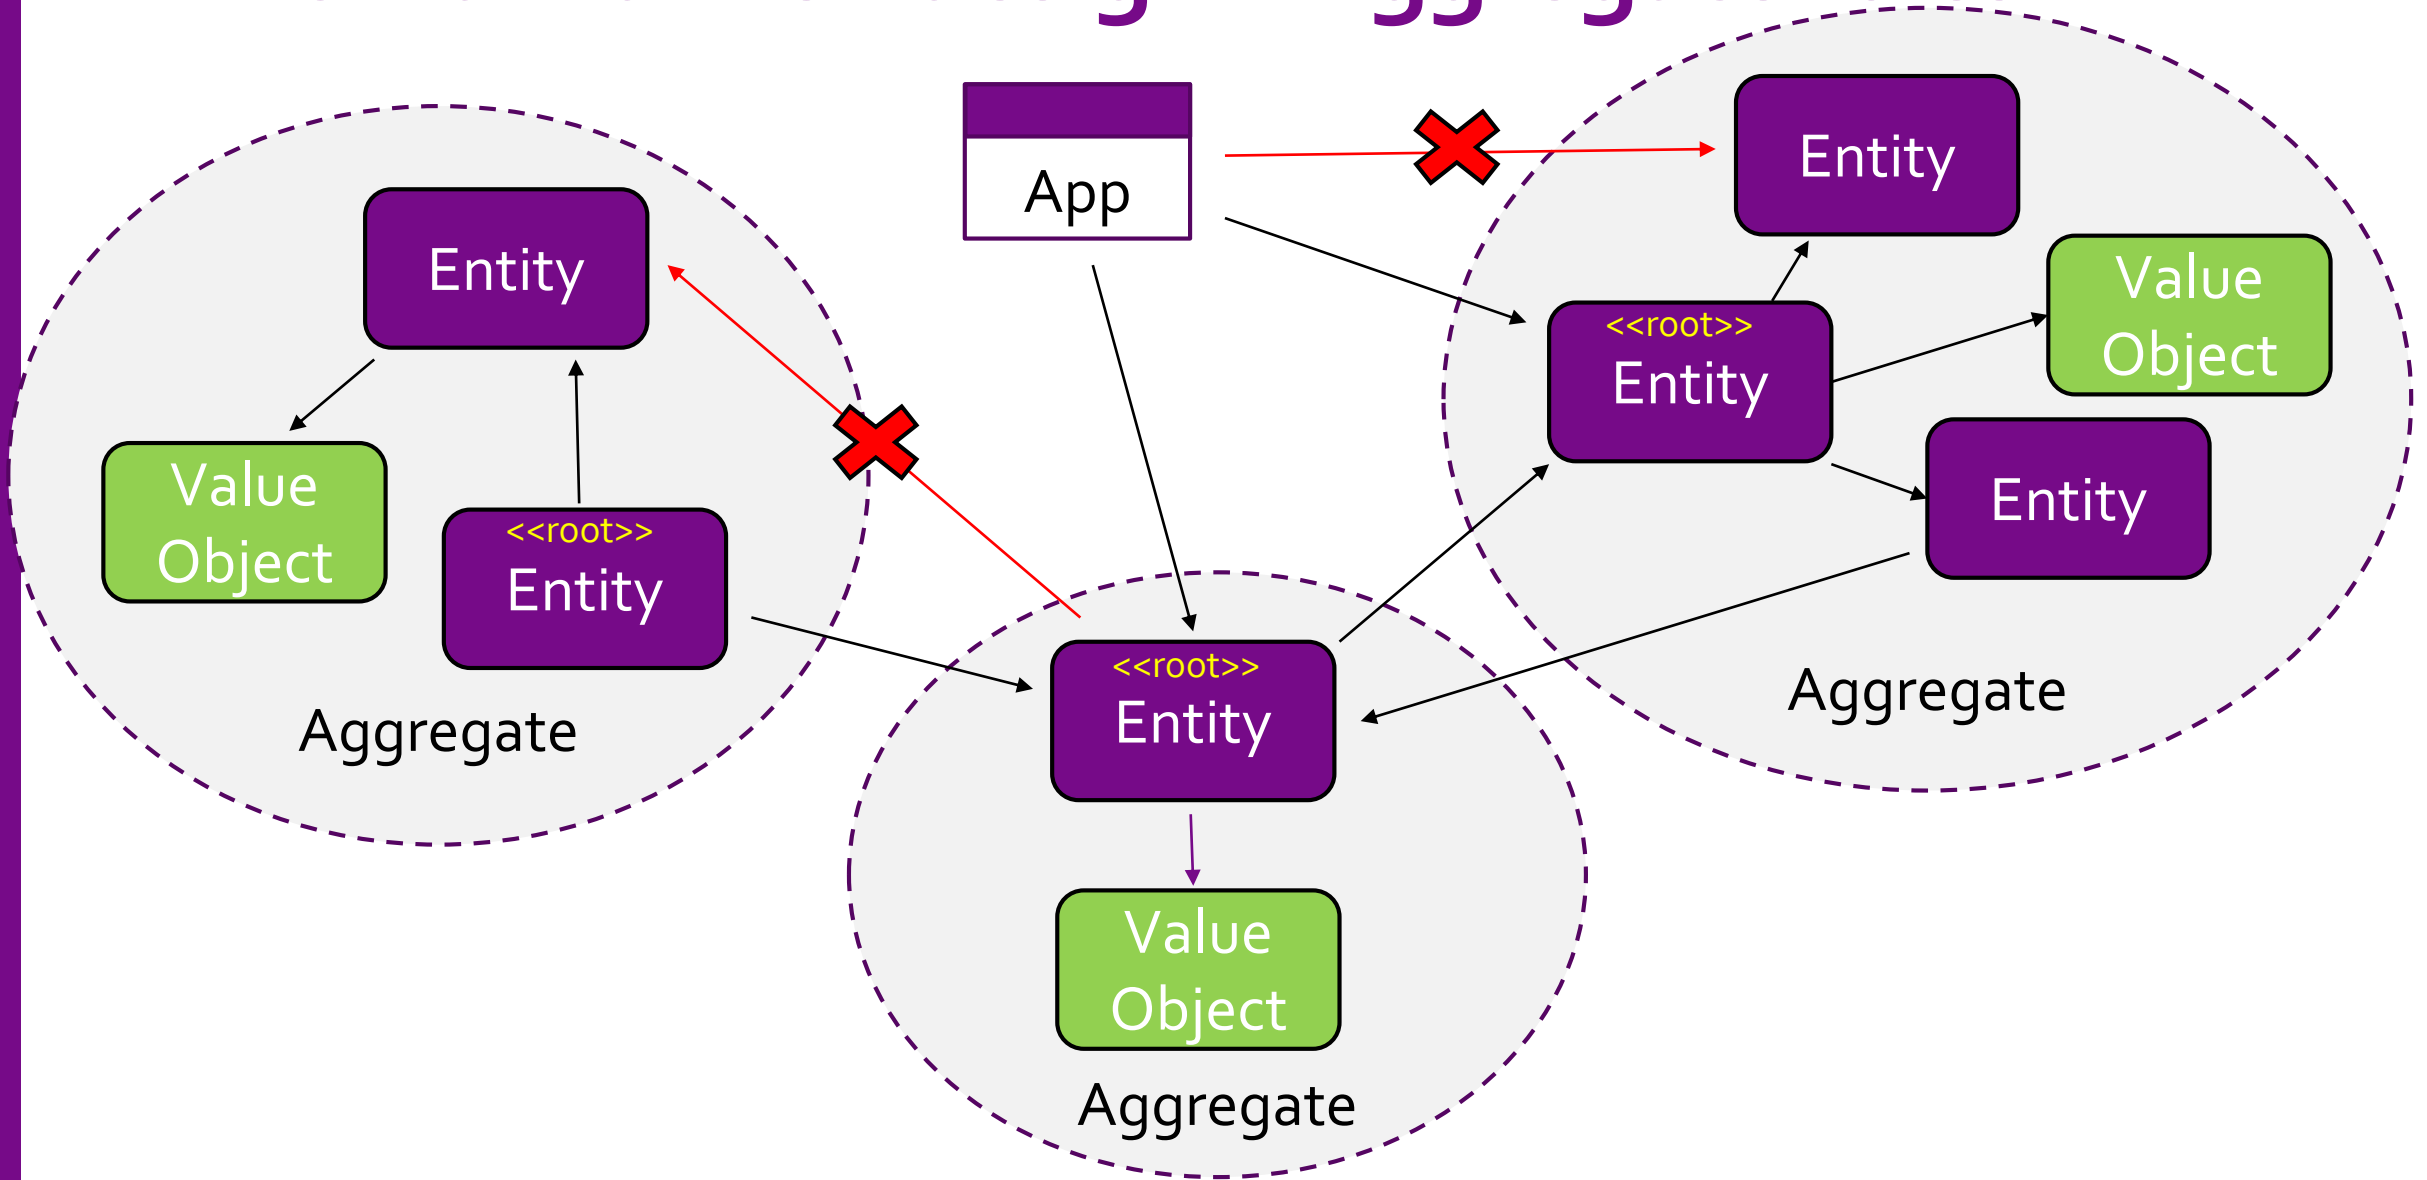
# Domain driven design - Aggregate

- Domain objects are divided into *Entities* and *Value objects*
  - Entities have an identity
    - Examples: Customer, Product, Contract, Car, ...
  - Value objects don't have an identity and are immutable
    - Able to identify based on the value of its attributes
    - Examples: Address, Temperature, Amount (\$), ...

# Domain driven design - Aggregates

- Group entities and value objects into *Aggregates*
  - Set of coherent objects with their relationships
- 1 object is the *Aggregate Root*
  - Only access-point for the aggregate
  - Ensures consistency for the entire aggregate
  - Has a single version# (optimistic concurrency control)

# Domain driven design – Aggregate rules



# Domain driven design

- Design your system around the business domain
- Do this together with subject matter experts (business analysts / end-users / ...)
- Speak the same language in documentation, design and code
  - “Ubiquitous Language”
- Divide the business into *bounded contexts*
  - Autonomous part of the business domain
  - Local definitions of business entities  
(naming based on the ubiquitous language)
  - Make technical choices per bounded context
    - Polyglot ‘X’ (best fit for purpose)
    - Architecture style (layered / onion / monolith / CQRS / ...)



distributed computing Onion architecture  
Event-based architecture

Domain modelling Actor model CQS

Web-scale architecture

Domain driven design CAP  
Distributed systems

Event streaming Serverless computing

Ubiquitous language

Web-scale IT

CQRS

Continuous delivery

Bulkhead

Partition tolerance

Consistency

NOSQL

Aggregate

CAP

Hexagonal architecture

Domain driven design

BASE

DevOps

Bounded context

Cloud computing

Circuit breaker

Event sourcing

Polyglot

X

DDD


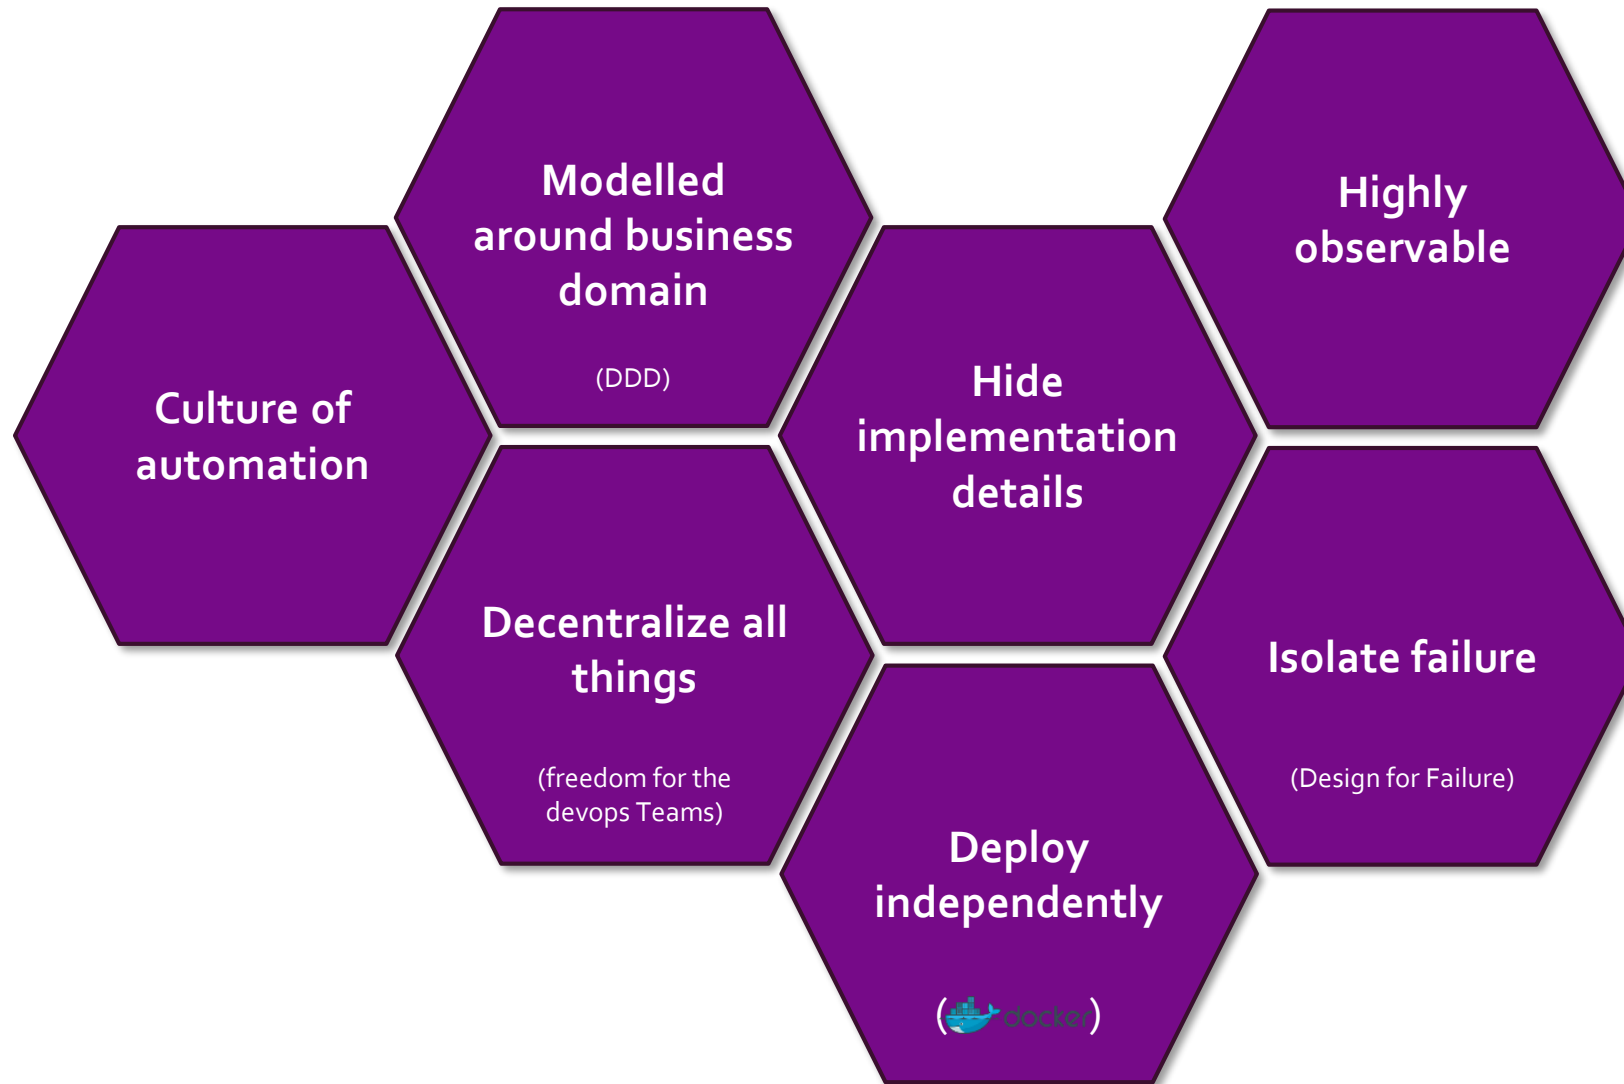
Design for failure

Eventual consistency

# Microservices

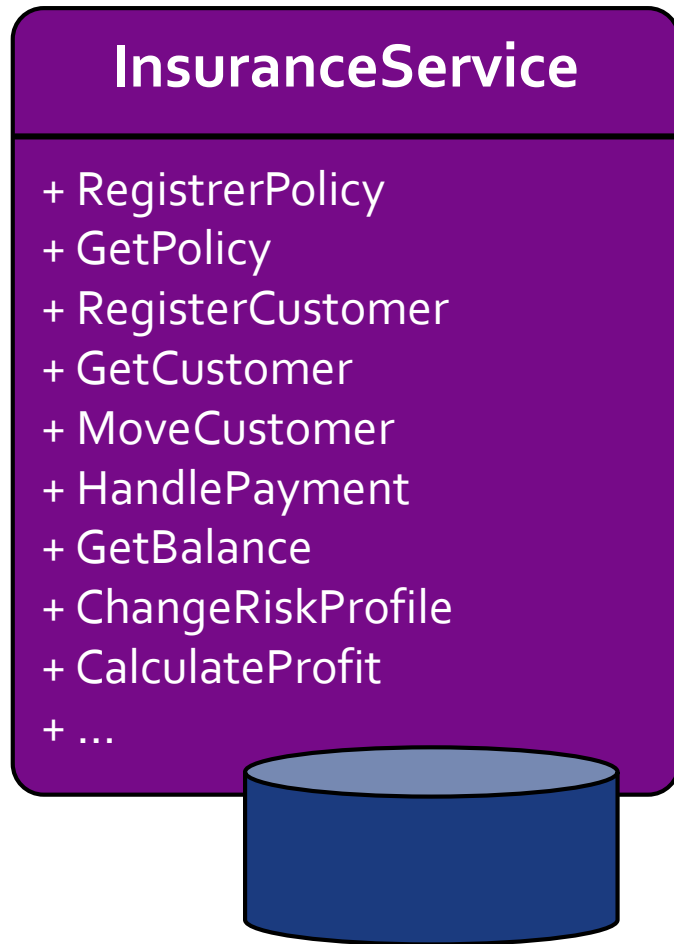
- “Small” autonomous services that cooperate
  - Designed around business domains /capabilities (DDD bounded contexts)
  - Simple to scale-out
  - High cohesion / low coupling
- A Microservice is specialized in 1 thing
  - Single responsibility principle

# Microservices principles

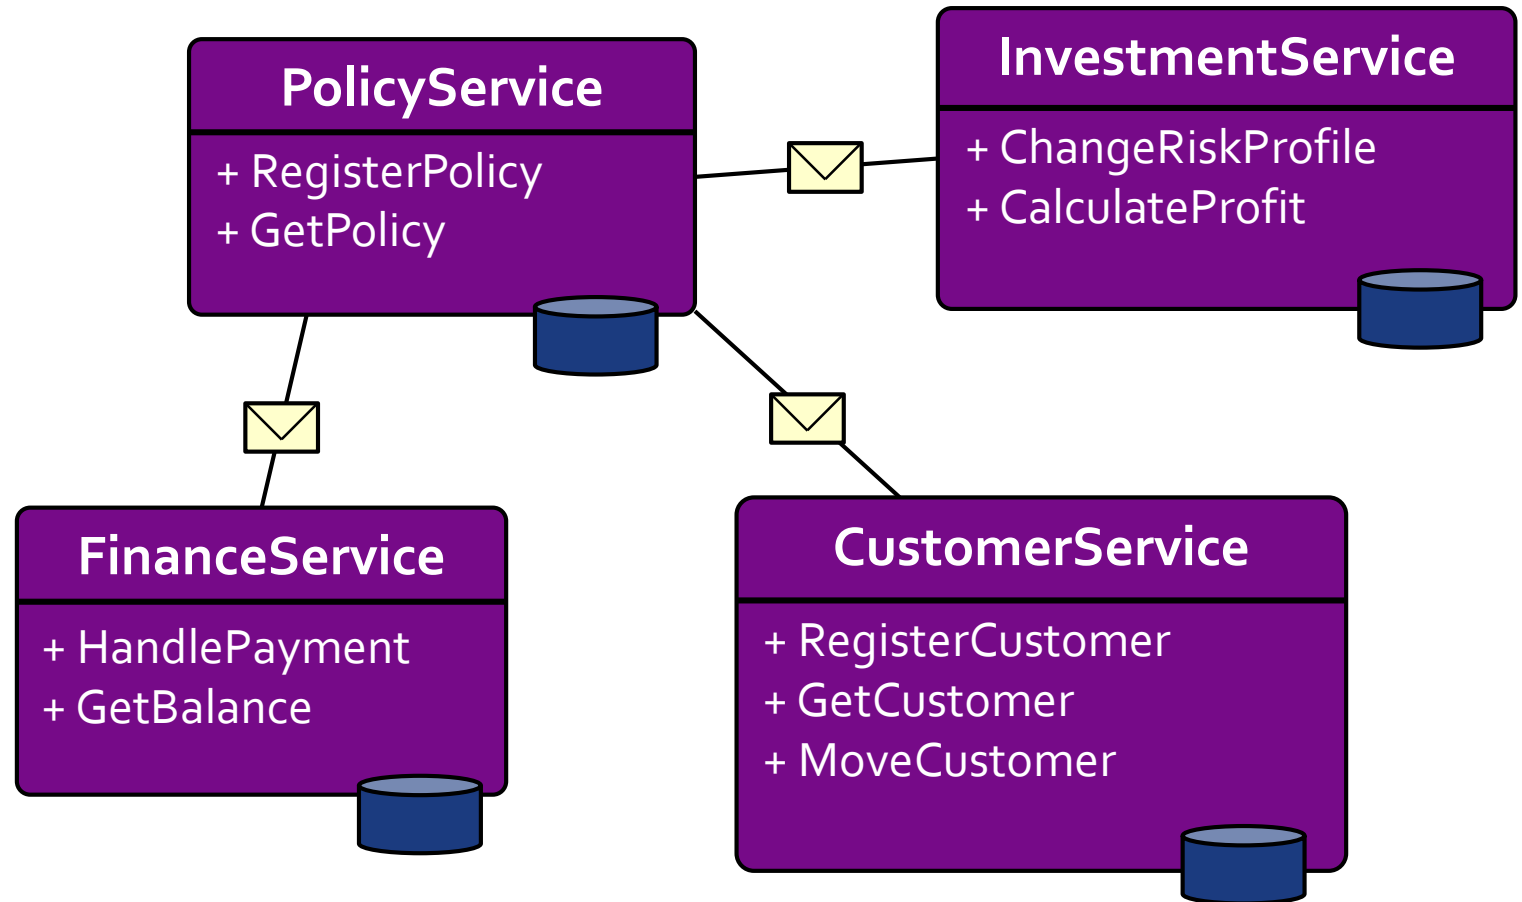


# Microservices

## Traditionally (SOA)



## Microservices



# Microservices

- Microservices communicate using “lightweight” protocols
  - HTTP (Rest API + JSON) / TCP + ProtoBuf / Own implementation
  - Choose between *open* or *fast*
- Primarily asynchronous communication


# Microservices

- Because of Microservice autonomy, certain technology-choices can be made per service:
  - **Architecture:** 3-tier | CQRS | Monolith | ...
  - **Language/runtime:** Java | C# | Scala | NodeJS | ...
  - **Persistence:** File System | MongoDB | Cassandra | ...
- This increases flexibility and makes sure you can choose the best solution per service (polyglot)

Event sourcing  
Design for failure  
distributed computing  
Micro services  
Web-scale IT  
Partition tolerance  
Event-based architecture  
NOSQL  
CQRS  
Polyglot X  
Agile  
CAP  
Ubiquitous language  
Serverless computing  
Bulkhead  
Distributed systems  
Consistency  
CQS  
Domain modelling  
BASE  
Event streaming  
DDD  
DevOps  
Onion architecture  
Aggregate  
Actor model  
Hexagonal architecture  
Circuit breaker  
CAP  
Continuous delivery  
Eventual consistency  
Domain driven design  
Bounded context  
Cloud computing

# Onion architecture


- Often we see a layered model being used
  - Presentation coupled to data-model
  - Business logic dependent on data-access and infrastructure layer
- Changes resonate throughout multiple layers





# Onion architecture

- A.k.a. "Hexagonal" architecture
  - Complexity and domain-logic lives in the center
  - Dependencies point "inwards" Domain-logic does not depend on infra / database



# Onion architecture

- Everything is decoupled using interfaces
  - Interfaces defined in the domain layer
  - No tight coupling with concrete implementation
  - Use of a dependency injection framework possible
- Unit-testing of business logic is simpler because of the lack of dependencies (easy mocking / stubbing)

Event sourcing  
Event-based architecture  
Domain driven design Consistency  
NOSQL  
CAP Web-scale IT Circuit breaker  
DDD Eventual consistency  
CQRS  
Actor model  
Bulkhead  
Event stroming  
Partition tolerance  
Bounded context BASE Ubiquitous language  
CQS Agile Polyglot X  
Aggregate Domain modelling  
CAP  
DevOps  
Distributed systems Serverless computing  
Continuous delivery  
Micro services Cloud computing  
Hexagonal architecture  
Onion architecture  
distributed computing  
Design for failure

# Actor model

- Pattern for building highly scalable distributed systems
- Particularly well suited for handling large concurrent workloads
- Several implementations exist:
  - Erlang
  - Scala.Akka
  - Akka.NET
  - Azure Service Fabric
  - ...

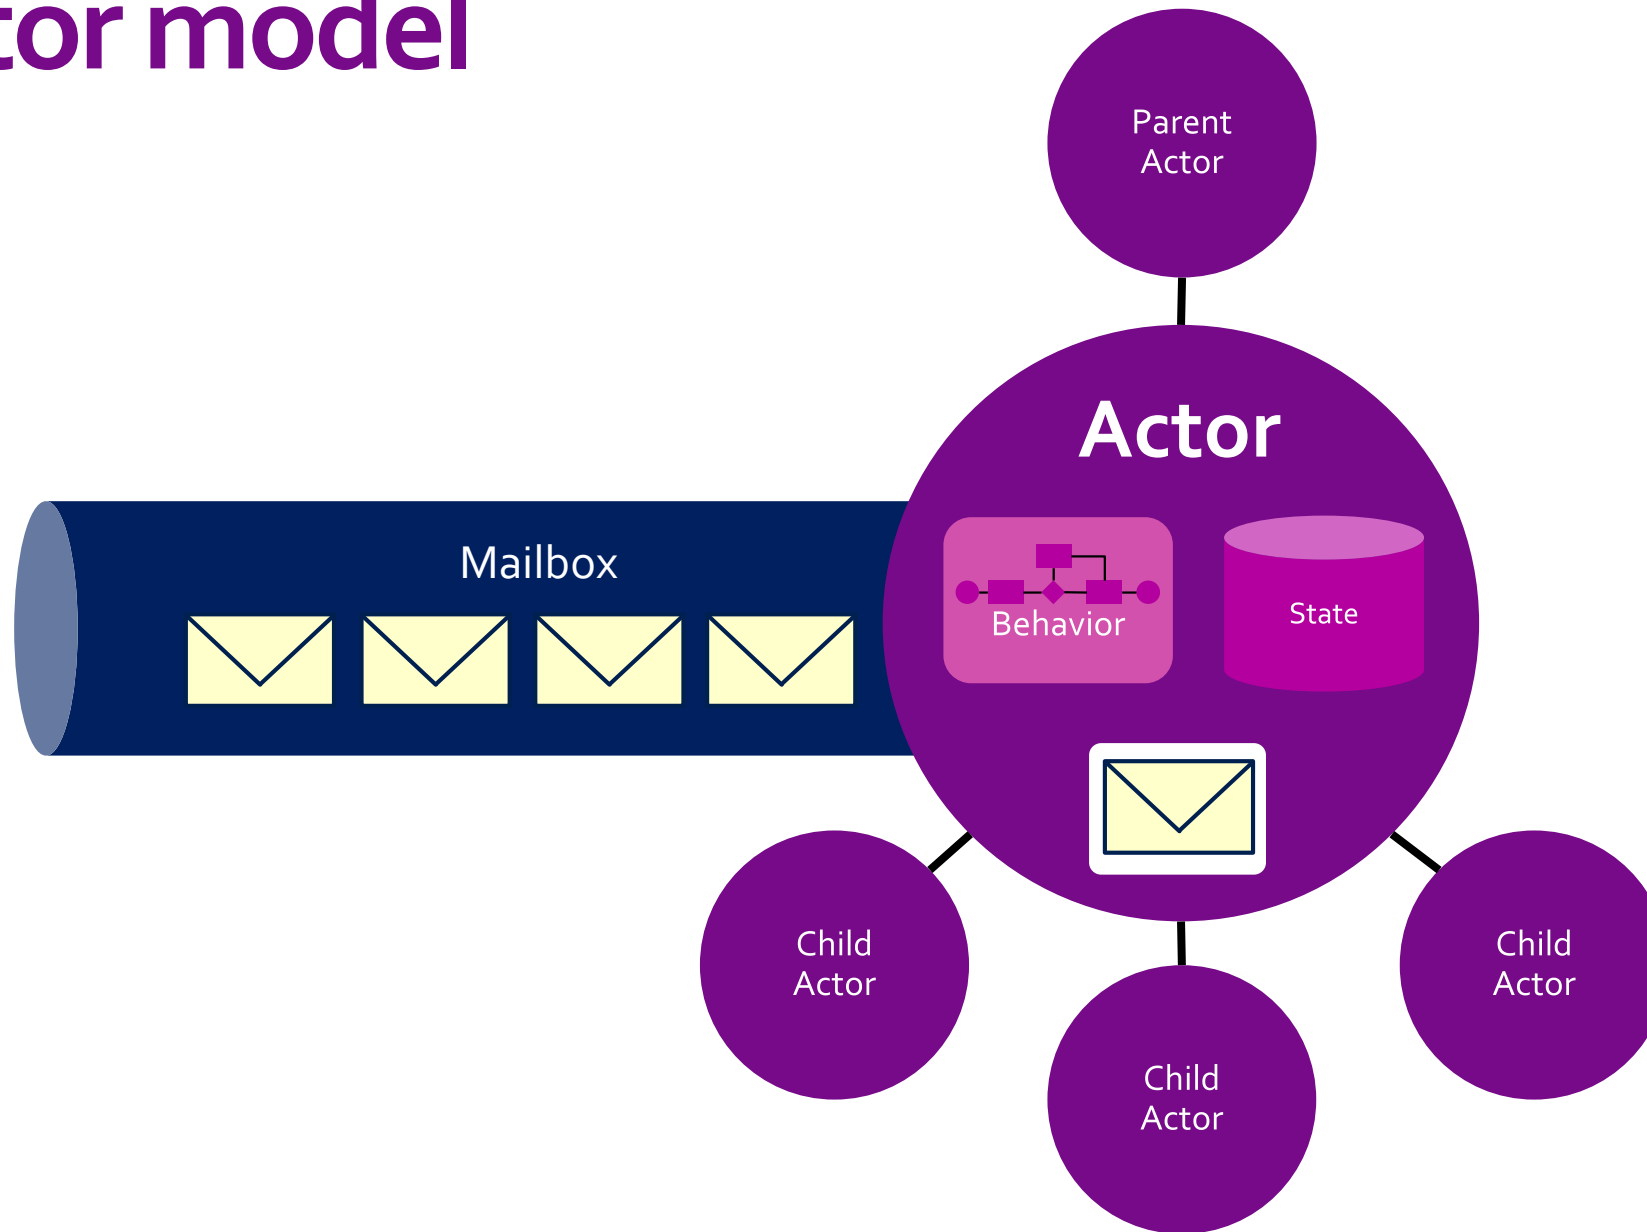
# Actor model

- Communication with an actor is done solely using asynchronous messages
- Messages arrive in the mailbox (queue) of the actor
- An actor handles exactly 1 message at a time
  - No concurrency issues when updating state

# Actor model

- An actor can handle a message in several ways
  - Execute certain (business-)logic
  - Update its own private state
  - Send messages to other actor(s)
  - Create child actors to offload work
  - Change its behavior
    - Like a finite state machine
    - Each state handles different message types

# Actor model



# Actor model - characteristics

- Actors are location-transparent
  - No difference in calling a local or a remote actor
  - Each actor has a unique address
- Actors are scalable
  - Scale-out by creating multiple actors
  - Clustering over multiple machines possible



# Actor model

- Not every problem can be solved using an Actor Model approach
- Sometimes looking at your problem differently can help
  - Partitioning of the work-load

# Thank you



**Prajeesh Prathap**

IT Architect



CONTACT

